

NewsLogic : AI-Powered Fake News Detection System

Course: Applied AI for Functional Leaders (IFSC 59903)

Team: Beta

Date: December 2025

1. Executive Summary

The AI-Powered Fake News Detection System is an applied artificial intelligence solution developed to evaluate the credibility of online text, news articles, and URLs by integrating multiple analytical models, rule-based detection mechanisms, and a large language model for contextual reasoning. It incorporates a transformer ensemble that includes a BERT-based classifier, a DistilBERT sentiment evaluator, optional Hugging Face Inference API fallback models, and a LLaMA-based explanation generator. Together, these components provide probabilistic predictions, contextual insights, and narrative explanations aligned with responsible AI principles emphasized in the Applied AI course.

The system uses a modern backend architecture built with FastAPI, Render, and Supabase to ensure secure authentication, automated deployment, structured logging, and user-specific history tracking. The Lovable/React frontend enables intuitive interaction, real-time analysis, and clean visualization of results. Environment-managed execution allows the system to use local models in development and high-availability cloud inference in production.

This project demonstrates the full AI lifecycle—from problem framing and data design to model implementation, governance, deployment, and monitoring—showcasing how functional leaders can safely and ethically deploy AI systems to address misinformation challenges.

2. Problem Landscape & System Objectives

2.1 The Misinformation Problem Landscape

Misinformation represents a significant societal and organizational challenge because digital platforms accelerate the circulation of unverified claims, emotionally charged narratives, and misleading content at a pace that far exceeds human review capacity. With social media algorithms prioritizing engagement-driven material, users are regularly exposed to sensational or polarizing posts that can distort perceptions, reinforce biases, and influence decision-making.

Artificial intelligence—particularly transformer-based models—is well suited for this problem because it can analyze large volumes of text rapidly, consistently, and without fatigue, enabling scalable credibility assessment across diverse sources. Transformers detect complex linguistic

cues, rhetorical framing techniques, contextual inconsistencies, and stylistic markers that frequently appear in fabricated or misleading information.

The Applied AI course emphasizes that systems combining automation and human oversight are essential in environments involving high-volume, repetitive analysis, making misinformation detection a strong applied AI use case. Organizations face reputational damage, compliance issues, operational risks, and financial consequences when misinformation spreads unchecked, reinforcing the need for robust and reliable detection mechanisms.

2.2 Why AI Is Well-Matched to the Problem

Transformer models such as BERT and DistilBERT are trained on large-scale datasets that teach them to understand contextual relationships, semantics, emotional tone, and linguistic patterns that differentiate truthful writing from deceptive claims. These models excel at identifying subtle indicators of misinformation, such as exaggerated rhetoric, emotional manipulation, unsupported generalizations, and shifts in narrative coherence.

Larger language models further strengthen this capability by generating structured, human-readable explanations that help users understand the reasoning behind model predictions. This reduces the “black box” effect by providing transparency into how the system interprets text, including the rhetorical patterns or contextual cues that influence the credibility assessment.

By using a combination of local model inference, rule-based logic, and fallback pathways through the Hugging Face Inference API, the system ensures robustness, continuity of service, and resilience to operational disruptions. This multimodel approach enables more comprehensive semantic understanding, credibility scoring, anomaly detection, and interpretive reasoning aligned with the complexities of modern misinformation.

2.3 System Objectives

- **Automate Credibility Assessment** : The system aims to automate the detection of misinformation by integrating transformer-based classifiers, sentiment models, and rule-based detectors to assess whether textual content is credible, unreliable, or deceptive.
- **Provide Transparent, Explainable Outputs** : The platform delivers clear explanations, including model confidence scores, linguistic patterns, sentiment indicators, and structured reasoning generated by the LLaMA explanation model, enabling users to understand how and why the system arrived at its conclusions.
- **Integrate Responsible AI Principles** : The system incorporates fairness checks, privacy protections, model transparency, audit logging, and human oversight mechanisms to ensure that the AI operates within ethical, secure, and accountable boundaries.

- **Demonstrate Deployment-Ready Architecture** : The system showcases a fully functional production pipeline with secure authentication, user history tracking, CI/CD deployment through Render, Supabase integration, inference monitoring, and modular code design suitable for real-world extension.
- **Enable Future Extensions** : The architecture is intentionally modular so that future enhancements—such as fact verification, multimodal analysis, and advanced model integration—can be incorporated without restructuring the entire system.

3. Stakeholders & Real-World Impact

3.1 Key Stakeholder Groups

The system benefits journalists who require reliable and rapid credibility assessments, researchers studying misinformation patterns, educators enhancing digital literacy, and trust-and-safety teams monitoring harmful or deceptive content. General users also gain a practical tool that improves their ability to evaluate and question online claims.

3.2 Impact on Digital Literacy and Public Understanding

Structured explanations and transparent scoring improve digital literacy by helping users recognize linguistic and rhetorical cues that frequently appear in misleading or manipulative content, fostering more informed online behavior.

3.3 Support for Organizational Integrity and Compliance

Organizations can integrate the system to evaluate content before dissemination, reducing reputational risks and supporting compliance with emerging misinformation-related guidelines and regulatory standards.

3.4 Contribution to a Healthier Information Ecosystem

By providing scalable, transparent, and interpretable credibility assessments, the system strengthens the information ecosystem, reduces exposure to misleading content, and promotes trust in digital communication environments.

4. System Overview

4.1 User Experience Flow

- **Secure Login and Session Establishment**

Users authenticate through Supabase, which issues a secure JWT token and enforces Row-Level Security. This ensures that every action—analysis, history retrieval, or data access—is tied to a verified user identity and protected under strict privacy controls.

- **Submit Text or URL for Analysis**

After logging in, users can paste a news article URL or enter raw text directly into the interface. The system automatically detects the input type; URLs undergo normalization, HTML parsing, and article extraction, while plain text moves straight to preprocessing.

- **Preprocessing and Inference Pipeline Execution**

The backend transforms the content through cleaning, normalization, segmentation for long inputs, and tokenization. The sanitized text is then processed through multiple components, including the BERT credibility classifier, DistilBERT/Twitter-RoBERTa sentiment model, rule-based detectors, and source reputation analysis.

- **Explanation Generation and Ensemble Decision Making**

The LLaMA-based explanation module (via Groq) generates a narrative summary that highlights linguistic patterns, rhetorical cues, and contextual signals influencing the classification. The system's ensemble logic fuses all model outputs and rule-based insights into a single credibility label, confidence score, and set of risk indicators.

- **Result Display and Secure History Logging**

The Lovable frontend presents the final output—classification label, confidence, explanation, and flags—through an intuitive interface. At the same time, all results are securely stored in Supabase under the user's account, enabling personalized history tracking and supporting responsible AI governance.

4.2 System Architecture Flow (With Render + GitHub + Lovable Integration)

Frontend Layer (Lovable)

The Lovable frontend serves as the primary user interface and manages all user interactions, including authentication, input submission, and results visualization. The frontend communicates with the backend through API calls directed at the FastAPI service hosted on Render. Each request includes a Supabase-issued authentication token to ensure secure access to inference endpoints and user-specific data.

Backend Processing Layer (FastAPI on Render)

The backend orchestrates URL extraction, text preprocessing, model inference, ensemble integration, rule-based detection, fallback execution, and logging. The backend is deployed on Render using continuous integration:

- The complete codebase resides in GitHub.
- Render automatically pulls updates from the GitHub repository whenever changes are pushed.

- Render builds the environment, installs dependencies, configures environment variables, and launches the FastAPI server using Uvicorn.

This automated GitHub → Render CI/CD pipeline ensures consistent deployments and reproducible environments. Once deployed, the backend exposes an HTTPS endpoint that the Lovable frontend uses to send text or URL analysis requests.

AI Model Stack

The core AI model stack includes BERT for credibility classification, DistilBERT for sentiment and emotional tone evaluation, a LLaMA-based reasoning model for narrative explanation, and rule-based detectors for rhetorical red flags such as exaggeration, emotionally charged language, and unsupported claims. A weighted ensemble mechanism fuses the outputs of these components into a unified credibility score and a human-readable explanation.

Database, Security & Governance Layer

- **Secure Authentication & Identity Controls (MFA, JWT, RLS)**
Supabase manages user identity with secure JWT tokens, optional multi-factor authentication, and strict Row-Level Security policies so each user can only access their own prediction history and logs.
- **Encryption for All Sensitive Data (At Rest & In Transit)**
All stored inputs, model predictions, metadata, and explanations are encrypted at rest, while TLS 1.2+ protects all data in transit; environment secrets such as API keys are isolated in Supabase's encrypted secret manager.
- **Data-Leak Prevention via Role Permissions & Policy-Based Access**
Least-privilege service roles, restricted API scopes, and signed storage URLs prevent unauthorized access or cross-user data exposure, ensuring article content and reasoning outputs remain private to each session.
- **Comprehensive Logging, Traceability & Audit Support**
Every prediction event is logged with user ID, timestamps, model versions, and explanation outputs, while Supabase audit logs track authentication events, schema changes, and backend actions to support responsible-AI governance.
- **Reliability, Backups, Monitoring & Automated Recovery**
Automated backups, restore capability, real-time monitoring, and anomaly detection protect the system from data loss, corruption, misuse, or operational failures, ensuring stable and trustworthy deployment.




Diagram : System Architecture

4.3 Frontend → Backend Communication

The Lovable application connects to the Render-hosted backend by making authenticated POST requests to the backend's public HTTPS endpoint (e.g., <https://service-name.onrender.com/analyze>) . CORS settings on the backend ensure secure cross-origin communication, and the system enforces authentication on every request. The frontend then receives structured JSON responses containing credibility scores, explanations, and metadata, which it formats into user-friendly visual components.

4.4 Render CI/CD Deployment Workflow

Render automates deployment through a multi-step process where it tracks changes in the GitHub repository, pulls the latest code, installs dependencies, applies environment variables (including Supabase and Hugging Face keys), builds the FastAPI service, and exposes the backend through a secure public URL. This workflow eliminates manual deployment steps, reduces human error, and ensures synchronized system updates.

5. Technical Components

5.1 AI Model Stack

- The system feeds the preprocessed input into multiple parallel models: a BERT classifier for credibility labeling, a DistilBERT sentiment model for emotional tone analysis, a rule-based detector for rhetorical red flags, and a URL/source analysis module for domain-level credibility signals.
- Outputs from all components are merged through an ensemble fusion and calibration layer that performs weighted scoring, consistency checks, and confidence thresholding to generate a single, unified credibility decision.
- In parallel, the LLaMA explanation model uses the processed input and model outputs to produce a grounded, human-readable justification that explains why the system classified the content as reliable, unreliable, or uncertain.




Diagram : AI Model Stack

5.2 Preprocessing & Data Pipeline

- User Login & Signup (Supabase Auth):** Users authenticate through Supabase using email/password or OAuth, and Supabase issues a secure JWT token. This token is attached to every API request, allowing the system to validate identity, enforce Row-Level Security, and associate all actions with the correct user account.
- Request Validation & Input Handling:** Once authenticated, the system detects whether the user submitted text or a URL. URLs are sanitized and securely fetched, while all inputs undergo normalization, cleaning, segmentation, and tokenization to create a consistent Processed Input for all downstream models.
- Unified Preprocessing Pipeline:** The pipeline extracts article content, removes scripts and noise, and standardizes text using structured rules. Both text and URL inputs feed into the same normalized representation, ensuring accurate and consistent performance across the entire model ensemble.
- Model Pipeline Execution:** The processed input branches into the BERT classifier, DistilBERT sentiment model, rule-based URL analyzer, and LLaMA explanation generator. Each model produces complementary outputs—credibility score, sentiment context, source analysis, and grounded narrative explanations.
- Supabase Logging & User History Storage:** After results are generated, the system securely stores credibility scores, model predictions, explanations, timestamps, and model

versions in Supabase under strict RLS policies. This builds a **personalized analysis history** for each user, enabling safe retrieval, traceability, and responsible-AI auditability.




Diagram : Data Preprocess Pipeline

5.3 Backend Infrastructure

FastAPI powers the backend with modular pipelines for inference, preprocessing, rule detection, and logging. In development, local models run directly; in production, Hugging Face Inference API is used for scalability. Render deploys the backend through automated CI/CD pipelines with secure environment management.

5.4 Explainability Layer

The explainability layer identifies influential words, sentiment indicators, and rhetorical signals. LLaMA provides coherent explanations that summarize model reasoning, highlight anomalies, and show how each component contributed to the final prediction.

5.5 Model Selection Rationale, Alternatives, and Prototype Constraints

5.5.1 Why These Models Were Used

- The system uses lightweight, free Hugging Face transformers (BERT-tiny for fake-news classification and DistilBERT/Twitter-RoBERTa for sentiment) because they require no GPU, run efficiently on Render’s free tier, and integrate smoothly with FastAPI and Supabase, allowing rapid end-to-end development.
- These models offer good baseline performance for credibility scoring while keeping compute and memory usage extremely low, which is essential in a constrained academic environment.
- For explanation generation, the project incorporates LLaMA-4-Scout-17B-16E-Instruct via the Groq API, providing fast, low-latency inference and high-quality reasoning without managing GPUs or large local models.
- The Groq API enables the prototype to deliver rich, human-readable narrative explanations, improving transparency and user trust while staying cost-effective.

5.5.2 What Other Models Could Perform Better

- Higher-capacity classifiers such as DeBERTa-v3 Large, RoBERTa-Large, ELECTRA-Large, or specialized fake-news models would improve accuracy, robustness, and contextual understanding.
- Long-document transformers (Longformer, BigBird, LED) would handle extended articles, transcripts, and multi-page content more effectively.
- Advanced reasoning models like GPT-4, GPT-4o, Claude 3, and LLaMA-3 70B would produce more coherent, well-grounded explanations and better detect subtle misinformation patterns.
- Retrieval-augmented architectures (RAG, FiD, T5-based fact verifiers) could enhance factual reliability by cross-checking claims against external knowledge sources or live databases.

5.5.3 Limitations of the Selected Models

- The current small models have limited context windows, which restricts their ability to understand long or complex articles.
- They perform classification based on patterns, not real evidence, meaning they cannot fact-check claims or verify sources automatically.
- Even the large LLaMA-4-Scout-17B model can hallucinate explanations, producing confident statements that may not be fully accurate.
- Satire, sarcasm, ideological framing, and subtle manipulation remain challenging for lightweight classifiers not trained specifically on misinformation datasets.
- These models are less robust than enterprise-grade systems and lack advanced components such as multi-modal reasoning, temporal tracking, or citation-level verification.

5.5.4 Why This Approach Works for a Prototype

- The selected models enable fast, inexpensive prototyping without requiring GPU infrastructure, making them ideal for a graduate-level applied AI project.
- They provide a complete demonstration of the AI lifecycle—data ingestion, preprocessing, ensemble scoring, explanation generation, model orchestration, and governance logic.
- The combination of lightweight Hugging Face models (efficient classification) and Groq-accelerated LLaMA explanations (strong interpretability) showcases a practical and balanced system design.
- The architecture highlights responsible AI principles—risk flags, source reputation checking, explainability, and ensemble decision logic—within a deployable full-stack application.
- This setup offers a clear pathway for future improvements, showing how the system can scale to more powerful models once compute and budget constraints are lifted.

6. Model Evaluation & Performance

6.1 Evaluation Dataset and Methodology

The system was evaluated using a dataset of 71 labeled news items exported directly from the deployed application’s analysis history. These entries reflect real user interactions and include a natural mix of Real and Fake news, along with examples of incomplete text, noisy content, and varying article lengths. This real-world dataset provides a practical basis for evaluating the prototype’s ability to classify credibility under realistic usage conditions, rather than a perfectly curated or balanced academic dataset.

Ground truth labels (“Real” or “Fake”) were mapped into binary form (Real = 1, Fake = 0), and model predictions (“REAL” or “FAKE”) were mapped into the same structure. Evaluation was conducted in a controlled Colab environment using standard scikit-learn metrics for classification

analysis. Because the system logs a single confidence score representing certainty in its predicted label, these values were converted into a probability of the Real class to compute meaningful ROC curve and AUC metrics. The evaluation therefore reflects both classification correctness and the model’s ability to discriminate between Real and Fake content across varying probability thresholds.

6.2 Quantitative Results

The system achieved an accuracy of 83.10%, demonstrating strong overall correctness across a diverse range of user-submitted content. Precision for the Real class was 0.75, meaning that when the system classified an article as Real, it was correct 75% of the time. The recall score of 0.90 shows that the system successfully identifies the majority of Real news articles, and the resulting F1 score of 0.818 reflects a balanced relationship between precision and recall. After converting logged confidence scores into a proper probability of the Real class, the system achieved a ROC AUC of approximately 0.91, indicating excellent discriminatory capability between Real and Fake articles.




Diagram : Confusion Matrix

The confusion matrix further illustrates model behavior, with 32 true negatives, 27 true positives, 9 false positives, and 3 false negatives. Most errors occurred when Fake articles were incorrectly classified as Real, which is consistent with the relatively lower precision score. Overall, these

results confirm that the model ensemble performs reliably and is well-suited for a functional prototype deployed in real-world conditions.

6.3 Interpretation of ROC AUC

The ROC curve illustrates the trade-off between the True Positive Rate and False Positive Rate across multiple classification thresholds. AUC (Area Under the Curve) provides a threshold-independent measure of model quality. The system's **AUC of ~0.91** indicates that in approximately 91% of randomly selected Real-versus-Fake pairs, the model assigns a higher credibility




Diagram : ROC Curve

probability to the Real article. This is a strong indicator of effective ranking and separation between the two classes, even when threshold-based metrics may show asymmetry. The high AUC demonstrates that the ensemble—despite using lightweight and free transformer models—successfully captures meaningful linguistic and contextual cues associated with credible writing.

6.4 Error Patterns and Practical Implications

An analysis of misclassified cases reveals several consistent patterns. First, Fake articles that employ neutral or minimally sensational language were sometimes misinterpreted as Real, reflecting the limitations of stylistic classification without factual grounding. Second, articles with very short or incomplete content provided insufficient context for reliable classification, leading to mislabeling. Third, HTML extraction inconsistencies—such as missing metadata or disrupted article structure—occasionally affected the available text during preprocessing, contributing to ambiguous predictions.

Despite these edge cases, the high recall for Real news and strong AUC score demonstrate that the system is effective at distinguishing credible content from fabricated or manipulated writing. The observed error patterns highlight opportunities for future enhancement, such as integrating fact-verification mechanisms, improving long-text handling, or incorporating richer source credibility signals. In practical terms, the system performs well as a decision-support tool that flags potential misinformation, provides explainable reasoning, and assists users in making informed judgments, while still maintaining the need for human oversight.

6.5 Case Study Analysis

Case Study 1 — Misclassified by the Fallback Model

Reuters published a series of verified Twitter statements made by former U.S. President Donald Trump. The article includes policy commentary, political accusations, and strongly-worded language (“Crooked Hillary,” “Dossier is bogus,” “FBI CANNOT VERIFY CLAIMS”). The outlet clearly states that the statements are unedited and the opinions expressed are Trump’s.

Ground Truth: Real

Primary Model (BERT): Real (99.9%)

Fallback Model (DistilBERT): Fake (99.7%)

Final Ensemble: Real

Explanation and Analysis

- **DistilBERT is triggered by emotional political language:** The fallback model associates highly charged political rhetoric with misinformation because similar patterns are common in fake-news training datasets. As a result, it produces a near-certain Fake prediction even though the content is accurately quoted.
- **DistilBERT may only see a truncated portion of the text:** When text exceeds its token limit, the model may not receive the Reuters attribution or explanatory framing. Without this context, the input resembles a raw political rant, increasing the likelihood of a false fake classification.

- **BERT sees more context and handles news-style reporting better** : The primary model processes longer inputs (up to ~512 tokens) and learns structural cues of formal journalism. It correctly interprets the article as a legitimate report quoting real statements.
- **DistilBERT is biased toward labeling political content as fake** : Because the fallback model was trained on datasets containing a high proportion of political misinformation, it developed a bias toward classifying emotionally intense political statements as fake—even when they appear in real articles.
- **The ensemble behaves correctly because BERT has a higher weight** : The final system output is Real because the ensemble assigns more influence to BERT (0.70). This weighting ensures that contextual understanding dominates stylistic misinterpretations, preventing a misclassification.

Case Study 2 —Misclassified by the Primary Model

A fabricated “Digital Access Tax” memo claims that the U.S. government will impose monthly fees on digital platform usage starting in 2026. The memo has no official source, and no reputable outlet has confirmed its legitimacy. The text is written in professional journalistic style and includes neutral disclaimers such as “not been verified” and “the Treasury Department has not responded.”

Ground Truth: Fake

Primary Model (BERT): Real (99.9%)

Fallback Model (DistilBERT): Fake (99.8%)

Final Ensemble: Real

Explanation and Analysis

- **BERT is strongly influenced by professional, neutral writing style** : Because the text mimics AP/Reuters-style reporting, BERT interprets the article as legitimate regardless of the factual accuracy of the claims. Writing quality, structure, and tone strongly bias BERT toward Real.
- **Neutral disclaimers make the text appear like legitimate fact-checking** : Phrases such as “document has not been verified,” “circulating online,” and “no legitimate outlet has confirmed” resemble credible reporting practices. BERT mistakenly treats these as signs of authenticity rather than indicators of misinformation.
- **The claim is subtle and bureaucratic, not sensational** : Transformer-based classifiers are more sensitive to extreme or emotionally charged language. This fabricated policy announcement is understated, making it difficult for BERT to detect as fake.

- **DistilBERT correctly identifies misinformation cues:** The fallback model recognizes the absence of reputable sourcing and the presence of common misinformation patterns (“unverified memo,” “circulating on social media”). These are strong signals of a fabricated narrative, producing an accurate Fake classification.
- **The ensemble still outputs Real because BERT has a higher weight:** Even though the fallback model correctly identifies the content as fake, the ensemble’s weighting prioritizes the primary model’s judgment. This leads to a final misclassification.

7. Data Handling Strategy & Operational Challenges

7.1 Secure Handling of User Data

All user inputs are linked to authenticated sessions through Supabase, ensuring that every prediction is traceable to the correct user while maintaining confidentiality. Prediction history is stored securely with row-level access controls, allowing each user to access only their own data. No personal identifiers beyond authentication tokens are stored, and all data is protected through encryption during transmission and at rest. This design supports privacy-by-default principles and protects the system from unauthorized access.

7.2 Handling Real-World Data Quality Issues

The system manages noisy, incomplete, or inconsistently formatted text by applying a rigorous preprocessing pipeline that standardizes and cleans content before inference. If URL extraction fails due to structured HTML inconsistencies or paywalls, fallback extraction methods attempt to retrieve the article content using alternative parsing strategies. When articles exceed model token limits, segmentation ensures that long-form content is still analyzable without overloading the transformer models.

7.3 Handling Adversarial or Manipulative Inputs

The system is designed to resist adversarial behaviors, such as exaggerated rhetoric intended to trigger sentiment-based manipulation or subtle paraphrasing designed to evade detection. Rule-based detectors identify sensational or emotionally charged writing, while ensemble modeling reduces reliance on any single model that may be more vulnerable to adversarial strategies. Strong input validation prevents malicious attempts such as code injection or malformed payloads. Suspicious inputs are logged to support future model improvement.

7.4 Operational Risks and Reliability Measures

Operational risks include inference delays, API rate limits, malformed URL structures, unexpected model failures, and cloud infrastructure limitations. Render introduces additional risks such as

service cold starts, dependency resolution failures, and redeployment errors triggered by GitHub commits. To mitigate these issues, the system implements retry-safe API requests, structured logging of errors, fallback inference pathways through the Hugging Face API, and health check procedures. The backend's modular design also allows problematic components to be replaced or updated independently without affecting the entire system.

7.5 Future Data Needs and Retraining Readiness

To ensure long-term system reliability and adaptiveness, the platform logs challenging examples for future training cycles, stores model version metadata, and captures user feedback regarding incorrect or ambiguous predictions. This data supports iterative refinement of the model pipeline and lays the foundation for more advanced MLOps workflows, including automated retraining, drift monitoring, and continuous integration of updated datasets.

8. Ethics & Governance

8.1 Fairness and Bias Mitigation

The system continuously monitors outputs across a diverse range of writing styles, linguistic structures, and content categories to identify potential sources of bias that may disproportionately affect specific groups or topics. When misclassification patterns emerge, they inform targeted updates to model prompts, rule-based detection logic, and underlying model configurations to ensure that the system behaves equitably across different forms of user-generated content.

8.2 Accountability Through Auditability

All analysis events performed by the system are captured in detailed Supabase audit logs, which record essential information such as user identifiers, model versions used during inference, timestamps of requests, classifier outputs, and explanation metadata, thereby enabling full traceability and accountability for every decision the system makes.

8.3 Transparency and Explainability

The system provides structured, human-readable explanations, confidence scores, and detector-trigger summaries that collectively ensure users can understand not only the final predictions but also the intermediate reasoning steps, reducing the perception of the system as a black-box model and supporting informed decision-making.

8.4 Privacy and Data Protection

Only the minimum necessary metadata required for system operation is collected and stored, and strong access-control mechanisms—combined with encryption of data both in transit and at rest—ensure that all user information is handled in alignment with privacy-by-design principles and protected against unauthorized access or misuse.

8.5 Human Oversight and Correctability

Although the system automates credibility detection, it is intentionally designed to complement, rather than replace, human judgment by enabling users to review predictions, identify errors, and submit corrective feedback, thereby ensuring that human oversight remains a key component of the decision-making loop.

8.6 Alignment with Governance Frameworks

The system aligns with globally recognized AI governance frameworks—including the NIST AI Risk Management Framework, OECD AI Principles, and the EU AI Act—by emphasizing transparency, fairness, accountability, privacy, and proportional oversight. These alignment practices ensure that the system’s operation remains consistent with emerging regulatory and ethical standards, supporting responsible deployment in both academic and real-world environments.

8.7 Security and Robustness Measures

Robust security measures are implemented throughout the system, including input sanitization to prevent injection attacks, protected environment variables to secure credentials, and encrypted communication channels to safeguard user data. The AI pipeline incorporates adversarial-resistant ensemble logic that reduces vulnerability to manipulative inputs designed to exploit model weaknesses. Continuous monitoring, API rate limiting, and fallback mechanisms further enhance resilience against system disruptions and malicious attempts.

9. Limitations

9.1 Text-Only Analysis

The current system analyzes only text-based content and is unable to evaluate images, videos, or multimodal misinformation, which limits its effectiveness in detecting increasingly common multimedia deception techniques such as deepfakes, edited images, or misleading visual-text combinations.

9.2 Transformer Context Window Limits

Because transformer models like BERT and DistilBERT have limited context windows, lengthy articles must be segmented into smaller portions before analysis, reducing the system’s ability to evaluate long-range narrative coherence and potentially causing fragmented interpretations across segments.

9.3 LLM Reasoning Limitations

Although the LLaMA explanation model provides helpful interpretive reasoning, its outputs are not fact-verified and may occasionally produce explanations that are plausible but not fully

grounded in external evidence. As a result, explanations should be viewed as interpretive insights rather than definitive statements of truth.

9.4 Computational Costs

Running multiple models in a single inference pipeline—combined with fallback mechanisms to external APIs—introduces computational overhead that can increase latency and reduce throughput, particularly in cloud-hosted environments or under high user demand.

9.5 Platform-Level Constraints

Supabase provides efficient authentication and storage for small to medium-scale applications, but it may not meet the performance, compliance, and redundancy requirements of enterprise-level systems without additional architectural enhancements or infrastructure support.

9.6 Vulnerability to Adversarial Manipulation

Sophisticated adversarial techniques, such as strategically paraphrased misinformation or subtle rhetorical manipulations, may still bypass detection because transformer models and rule-based systems can struggle to identify well-crafted deceptive writing that avoids common linguistic cues.

9.7 Lack of Factual Grounding

The system evaluates credibility based on linguistic features and contextual reasoning rather than verifying claims against external reliable sources. Without retrieval-based grounding, the system cannot determine whether a statement is objectively true—it can only assess whether the writing resembles credible or deceptive patterns.

10. Future Enhancements

10.1 Multimodal Analysis

Expanding the system to evaluate images, videos, and combined text-media inputs would allow the platform to detect a broader spectrum of misinformation, including manipulated photos, edited videos, and multimodal deceptive narratives commonly found on social platforms.

10.2 Retrieval-Augmented Fact Verification (RAG)

Integrating retrieval-based fact verification would enable the system to compare claims against authoritative external sources, producing explanations backed by citations and reducing reliance on purely stylistic or linguistic analysis. This approach would significantly improve factual grounding and reliability.

10.3 MLOps Integration

Incorporating complete MLOps workflows—including automated model retraining pipelines, drift detection systems, performance dashboards, and continuous integration/testing—would enhance long-term maintainability and allow the system to adapt to evolving misinformation patterns.

10.4 Scalable Deployment

Scaling the backend through containerization technologies such as Docker and orchestration frameworks like Kubernetes would improve deployment flexibility, fault tolerance, and horizontal scaling, enabling the system to support much higher user volumes in production environments.

10.5 Explainability Improvements

Integrating advanced interpretability techniques such as SHAP values, Integrated Gradients, and counterfactual reasoning tools would deepen user understanding of model behavior and offer more granular insights into the relationship between input features and credibility predictions.

10.6 Enhanced User Experience

Improving the user interface through multilingual support, richer visualizations, adaptive explanation formats, and interactive analysis tools would make the system more accessible, intuitive, and informative for a wide range of users with different backgrounds and needs.

10.7 Continuous Learning Through User Feedback

Introducing a structured feedback loop where users can label incorrect predictions or provide clarifying context would allow the system to evolve dynamically over time, supporting supervised updates and helping reduce recurring errors.

11. Alignment with Course Topics

11.1 Reflection of the Applied AI Lifecycle

The system mirrors the full AI lifecycle emphasized in the course, including the stages of problem framing, model selection, data preparation, iterative development, validation, deployment through cloud infrastructure, monitoring, and opportunities for ongoing improvement, demonstrating how these theoretical concepts translate into a practical, operational AI solution.

11.2 Integration of Responsible AI Principles

The system incorporates fairness, transparency, accountability, privacy protection, and human oversight, aligning directly with the Responsible AI frameworks taught in the course and illustrating how ethical considerations can be woven into every stage of the design and deployment pipeline.

11.3 Demonstration of Technical and Governance Competencies

By combining transformer-based models, explanation generation, secure authentication mechanisms, audit logging, and structured deployment workflows, the project demonstrates mastery of both technical competencies and governance-oriented requirements expected in an applied AI project.

11.4 Practical Application of Course Learnings

The project serves as a concrete example of how theoretical instruction in AI development, system architecture, data governance, and ethical oversight can be transformed into a functional tool that addresses a real-world societal challenge.

11.5 Project Reflection & Lessons Learned

- **Model practicality vs. ambition:** The project highlighted the importance of choosing models that balance performance with deployment constraints. Lightweight transformers and Groq-hosted LLaMA provided the ideal mix for a prototype.
- **Architecture matters as much as modeling:** A large portion of the work involved designing preprocessing, modular pipelines, secure storage, and deployment workflows—reinforcing the course theme that AI systems require much more than just a model.
- **Explainability is essential:** Adding the LLaMA explanation model significantly improved user trust and interpretability, showing how narrative explanations help bridge model decisions and human understanding.
- **Responsible AI must be embedded early:** Privacy controls, auditability, and row-level security shaped the entire design process rather than being added at the end.
- **Real-world complexity:** Testing showed that misinformation detection is difficult even with AI, and revealed limitations around satire, emotional content, and long articles—demonstrating why human oversight remains critical.

12. Conclusion

The AI-Powered Fake News Detection System demonstrates how modern AI techniques, thoughtful governance structures, and user-centered design principles can be combined to evaluate the credibility of online content. Its multi-model ensemble, explainability mechanisms, secure backend infrastructure, and compliance with responsible AI principles offer a transparent and reliable approach to misinformation detection. Although the system faces limitations—particularly around multimodal content, factual grounding, and computational efficiency—its modular design supports future expansion and enhancement. This project ultimately shows how functional leaders can apply AI thoughtfully and ethically to create scalable, meaningful solutions that strengthen information integrity and support human decision-making.

References

- [1] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of deep bidirectional transformers for language understanding,” *Proc. NAACL-HLT*, 2019. [Online]. Available: <https://arxiv.org/abs/1810.04805>
- [2] V. Sanh, L. Debut, J. Chaumond, and T. Wolf, “DistilBERT, a distilled version of BERT: Smaller, faster, cheaper and lighter,” 2019. [Online]. Available: <https://arxiv.org/abs/1910.01108>
- [3] H. Touvron *et al.*, “LLaMA: Open and efficient foundation language models,” 2023. [Online]. Available: <https://arxiv.org/abs/2302.13971>
- [4] S. Vosoughi, D. Roy, and S. Aral, “The spread of true and false news online,” *Science*, vol. 359, no. 6380, pp. 1146–1151, 2018, doi: 10.1126/science.aap9559.
- [5] X. Zhou and R. Zafarani, “A survey of fake news: Fundamental theories, detection methods, and opportunities,” *ACM Computing Surveys*, vol. 53, no. 5, pp. 1–40, 2020, doi: 10.1145/3395046.
- [6] National Institute of Standards and Technology, “Artificial Intelligence Risk Management Framework (AI RMF 1.0),” U.S. Dept. of Commerce, 2023. doi: 10.6028/NIST.AI.100-1.
- [7] OECD, “OECD Principles on Artificial Intelligence,” 2019. [Online]. Available: <https://oecd.ai/en/ai-principles>
- [8] European Parliament and Council, “EU Artificial Intelligence Act,” 2024. [Online]. Available: <https://artificialintelligenceact.eu>
- [9] Supabase, “Supabase Authentication, Storage & Database Documentation,” 2024. [Online]. Available: <https://supabase.com/docs>
- [10] Render, “Render Cloud Platform Documentation,” 2024. [Online]. Available: <https://render.com/docs>
- [11] Hugging Face, “Hugging Face Inference Endpoints & Model Hub Documentation,” 2024. [Online]. Available: <https://huggingface.co/docs>